

# Step 0: Load Packages

Before you run next chunk, please follow the instructions to install all packages we need:

Pre-requirements:

numpy, random, pickle, time, xgboost, PIL, gist, csv, FFTW

(1) Install numpy, random, pickle

```
$ pip install numpy
```

```
$ pip install random
```

```
$ pip install pickle
```

(2) Install FFTW

FFTW download: <http://www.fftw.org> (<http://www.fftw.org>)

Install instruction: [http://www.fftw.org/fftw3\\_doc/Installation-on-Unix.html](http://www.fftw.org/fftw3_doc/Installation-on-Unix.html)  
([http://www.fftw.org/fftw3\\_doc/Installation-on-Unix.html](http://www.fftw.org/fftw3_doc/Installation-on-Unix.html))

```
$ ./configure --enable-single --enable-shared
```

```
$ make
```

```
$ sudo make install
```

(3) Install gist

Download lear\_gist: <https://github.com/tuttieeee/lear-gist-python> (<https://github.com/tuttieeee/lear-gist-python>)

```
$ sudo python setup.py build_ext
```

```
$ python setup.py install
```

If fftw3f is installed in non-standard path (for example, HOME/local), use -I and -L options:

```
$ sudo python setup.py build_ext -I HOME/local/include -L HOME/local/lib
```

(4) Install xgboost

Install instructions:

[https://www.ibm.com/developerworks/community/blogs/jfp/entry/Installing\\_XGBoost\\_on\\_Mac\\_OSX?lang=en](https://www.ibm.com/developerworks/community/blogs/jfp/entry/Installing_XGBoost_on_Mac_OSX?lang=en)

([https://www.ibm.com/developerworks/community/blogs/jfp/entry/Installing\\_XGBoost\\_on\\_Mac\\_OSX?lang=en](https://www.ibm.com/developerworks/community/blogs/jfp/entry/Installing_XGBoost_on_Mac_OSX?lang=en)).

Problems you might encounter when you install: <https://stackoverflow.com/questions/36211018/clang-error-errorunsupported-option-fopenmp-on-mac-osx-el-capitan-buildin>

(<https://stackoverflow.com/questions/36211018/clang-error-errorunsupported-option-fopenmp-on-mac-osx-el-capitan-buildin>).

In [1]:

```
import GIST
import pandas as pd
import random
import pickle
import time
import xgboost
```

## Step 1: Read Test Pictures Information

Before you run next chunk, please make sure you meet following requirements:

- (1) Make sure path variable is where you store all your test images
- (2) Make sure 5000 SIFT feature descriptors of your test images are stored in the data folder as feature\_sift\_test.csv
- (3) Make sure label of your test images are stored in the data folder as label\_test.csv

In [2]:

```
path = "/Users/siyi/Documents/Study-Columbia/17FALL/GR5243-Applied-Data-Science/Project3/training_set/images2"
GIST.feature_output(path)
gist_new = pd.read_csv('feature.csv', skiprows=1, header = None).iloc[:, 1:]
sift_new = pd.read_csv('../data/feature_sift_test.csv').iloc[:, 1:]
label_new = pd.read_csv('../data/label_test.csv').iloc[:, 1]
feature = pd.concat([sift_new, gist_new], axis=1)
feature.columns = ['x' + str(i+1) for i in range(5000)] + ['f' + str(i+1) for i in range(960)]
```

## Step 2: XGBoost Model

In [3]:

```
# require X_test, y_test
X_test = feature
y_test = label_new
```

In [4]:

```
# load the baseline model
filename = '../output/model_baseline.sav'
xgb_1 = pickle.load(open(filename, 'rb'))

# load the tuned model
filename = '../output/model_tuned.sav'
xgb_2 = pickle.load(open(filename, 'rb'))
```

In [5]:

```
start = time.time()
print("Baseline: ")
pred = xgb_1.predict(X_test)
y_label = y_test.values
print ('classification error=%f' % (sum([pred[i] != y_label[i] for i in range(len(y_label))]) / float(len(y_label)) ))

start = time.time()
print("Tuned: ")
pred = xgb_2.predict(X_test)
y_label = y_test.values
print ('classification error=%f' % (sum([pred[i] != y_label[i] for i in range(len(y_label))]) / float(len(y_label)) ))
```

```
Baseline:
classification error=0.000000
Tuned:
classification error=0.000000
```

In [ ]:

# baseline\_main

Yufei Zhao

2017/10/29

```
print(R.version)
```

```
##  
## platform      _  
## arch          x86_64-mingw32  
## os            mingw32  
## system        x86_64, mingw32  
## status  
## major         3  
## minor         4.2  
## year          2017  
## month         09  
## day           28  
## svn rev       73368  
## language      R  
## version.string R version 3.4.2 (2017-09-28)  
## nickname      Short Summer
```

```
if(!require("gbm")){  
  install.packages("gbm")  
}
```

```
## Loading required package: gbm
```

```
## Loading required package: survival
```

```
## Loading required package: lattice
```

```
## Loading required package: splines
```

```
## Loading required package: parallel
```

```
## Loaded gbm 2.1.3
```

```
if(!require("dplyr")){  
  install.packages("dplyr")  
}
```

```
## Loading required package: dplyr
```

```
##  
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':  
##  
## filter, lag
```

```
## The following objects are masked from 'package:base':  
##  
## intersect, setdiff, setequal, union
```

```
library(gbm)  
library(dplyr)
```

## Step 0: set seed.

In order to obtain reproducible results, `set.seed()` whenever randomization is used.

```
set.seed(1)
```

## Step 1: set up controls for evaluation experiments.

In this chunk, we have a set of controls for the evaluation experiments.

- (T/F) cross-validation on the training set
- (number) K, the number of CV folds
- (T/F) process features for training set
- (T/F) run evaluation on an independent test set
- (T/F) process features for test set

```
run.cv=TRUE # run cross-validation on the training set  
K <- 3 # number of CV folds  
#run.feature.train=TRUE # process features for training set  
run.test=TRUE # run evaluation on an independent test set  
#run.feature.test=TRUE # process features for test set
```

Using cross-validation or independent test set evaluation, we compare the performance of different classifiers or classifiers with different specifications. In this example, we use GBM with different `depth`. In the following chunk, we list, in a vector, setups (in this case, `depth`) corresponding to models that we will compare. In your project, you maybe comparing very different classifiers. You can assign them numerical IDs and labels specific to your project.

```
model_values <- seq(3, 11, 4)
model_labels = paste("GBM with depth =", model_values)
```

## Read in data

```
# lables(0 for muffin, 1 for chicken, 2 for dog)
labels <- read.csv("../data/label_train.csv",header=TRUE)
colnames(labels)[2] <- "labels"
```

```
sift_data <- read.csv("../data/sift_train.csv",header=TRUE, stringsAsFactors = FALSE)
```

## divide into training and sampling data set

```
sampling <- sample(seq(1:3000),3000*0.25) %>% sort()
label_test <- labels[sampling,-1]
label_train <- labels[-sampling,-1]
dat_test <- sift_data[sampling,-1]
dat_train <- sift_data[-sampling,-1]
```

## load train and test method

```
source("../lib/baseline_train.R")
source("../lib/baseline_test.R")
```

## Model selection with cross-validation

- Do model selection by choosing among different values of training model parameters, that is, the interaction depth for GBM in this example.

```
# small range case !!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!
#dat_train <- sift_data[1:100,-1]
#label_train <- labels[1:100,-1]
#dat_test <- sift_data[101:200,-1]
#label_test <- labels[101:200,-1]
```

```

source("../lib/baseline_cv.R")

if(run.cv){
  err_cv <- array(dim=c(length(model_values), 2))
  for(k in 1:length(model_values)){
    cat("k=", k, "\n")
    err_cv[k,] <- cv.function(dat_train, label_train, model_values[k], K)
  }
  save(err_cv, file="../output/baseline_err_cv.RData")
}

```

```

## k= 1
## k= 2
## k= 3

```


Visualize cross-validation results.

```

if(run.cv){
  load("../output/baseline_err_cv.RData")
  #pdf("../fig/cv_results.pdf", width=7, height=5)
  plot(model_values, err_cv[,1], xlab="Interaction Depth", ylab="CV Error",
        main="Cross Validation Error", type="n", ylim=c(0, 0.6))
  points(model_values, err_cv[,1], col="blue", pch=16)
  lines(model_values, err_cv[,1], col="blue")
  arrows(model_values, err_cv[,1]-err_cv[,2], model_values, err_cv[,1]+err_cv[,2],
        length=0.1, angle=90, code=3)
  #dev.off()
}

```

## Cross Validation Error



- Choose the “best” parameter value

```
model_best=model_values[1]
if(run.cv){
  model_best <- model_values[which.min(err_cv[,1])]
}

par_best <- list(depth=model_best)
```

- Train the model with the entire training set using the selected model (model parameter) via cross-validation.

```
tm_train=NA
tm_train <- system.time(fit_train <- train(dat_train, label_train, par_best))
```

```
## Warning in gbm.fit(x = dat_train, y = label_train, n.trees = 1000,
## distribution = "multinomial", : variable 1012: V1012 has no variation.
```

```
## Warning in gbm.fit(x = dat_train, y = label_train, n.trees = 1000,
## distribution = "multinomial", : variable 2688: V2688 has no variation.
```



```
## Warning in gbm.perf(fit_gbm, method = "OOB", plot.it = FALSE): OOB
## generally underestimates the optimal number of iterations although
## predictive performance is reasonably competitive. Using cv.folds>0 when
## calling gbm usually results in improved predictive performance.
```

```
save(fit_train, file="../output/baseline_fit_train.RData")
```

## Step 5: Make prediction

Feed the final training model with the completely holdout testing data.

```
tm_test=NA
if(run.test){
  load(file="../output/baseline_fit_train.RData")
  tm_test <- system.time(pred_test <- test(fit_train, dat_test))
  save(pred_test, file="../output/baseline_pred_test.RData")
}
```

## Calculate the prediction rate

```
load(file="../output/baseline_pred_test.RData")
correct_rate <- mean(pred_test == label_test)
cat("Correct rate for prediction = ", correct_rate[1], "\n")
```

```
## Correct rate for prediction =  0.7546667
```

## Summarize Running Time

Prediction performance matters, so does the running times for constructing features and for training the model, especially when the computation resource is limited.

```
cat("Time for training model=", tm_train[1], "s \n")
```

```
## Time for training model= 2775.39 s
```

```
cat("Time for making prediction=", tm_test[1], "s \n")
```

```
## Time for making prediction= 0.11 s
```